

Amazon Books Model

Claudiu Reditu 266129

Stefan Harabagiu 266116

Supervisors: Ole Ildsgaard Hougaard

VIA University College

Software Engineering

Semester 6

29-03-2020

Model Description

The starting point for building our model is the conceptual model of the Amazon Books Database:

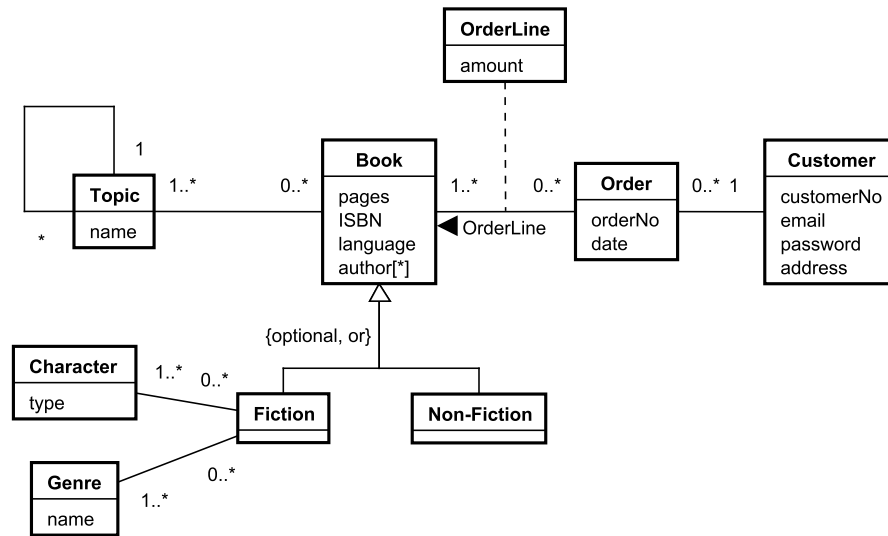


Figure 1: Conceptual Diagram

Decisions and Reasons

The first choice is deciding how Topic should handle the recursive relationship. The Parent Reference pattern is the simplest solution that helps in performing queries to find all topics.

The second choice is to remove entities, such as OrderLine, because they are no longer required to solve many to many relationships. They can be solved by storing an array of ids in one of the entities. All of these relationships are modelled in the same way.

The third choice is to have the following fields in book: type(Fiction, Non-Fiction), an array of objects for characters and an array to reference entries in genre. Genres are not book specific as most characters are. Books are genre specific, so the same genre appears frequently in many books. The array of objects for characters can work as long as there are not a lot of characters included. Usually it is important to have the few main characters that are specific to that book. This is the reason why characters are present in book.

The fourth choice is having authors just as an array of names.

The result is Fig 2. It describes the relationship between the collections in Mongo:

- The Parent Reference pattern provides a simple solution to tree storage, but requires multiple queries to retrieve subtrees
- The Child References pattern provides a suitable solution to tree storage as long as no operations on subtrees are necessary. This pattern may also provide a suitable solution for storing graphs where a node may have multiple parents.
- The Array of Ancestors pattern - no specific advantages unless you constantly need to get path to the node

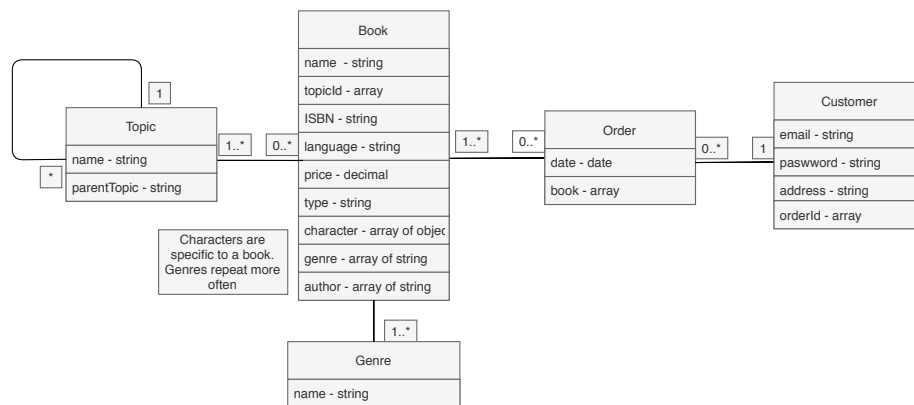


Figure 2: Amazon Books Mongo Diagram

Consequences

A consequence of the first choice is the easiness of using graphLookup to find each topic.

With the removal of entities that solves many to many cases, joining and performing different operations required more steps in the Mongo pipeline than it would otherwise in SQL.

Having type as a field instead of two different tables, while adding genre and characters when needed, makes it easy to handle the exercises in which it is mentioned. Less lookups are performed and the database is filled with less redundant data.

Usually, books don't have more than four authors, but one author can have many books. Having author just as an array of names makes it more taxing to perform queries, but it means less redundant data in the database. If more data is needed on an author than name, it can be a good idea to make it separate from book, but in this case it works well.

The Exercises

The most difficult exercises are the ones that require to keep the values even though their aggregates are 0. For example, is it difficult to solve exercise nine because all books have to be in the response, even the ones that aren't being bought. Without this condition, the exercise is straight-forward. There are difficulties in making sure the syntax is correct. Even with simple use cases the queries can seem overwhelming.

Otherwise, most exercises are straightforward and the way pipelines work makes it easy to handle the documents. Each step of the pipeline makes logical sense due to the structure of documents being based on the idea of objects from object-oriented programming languages. Aggregating in Mongo is easier than in SQL, if the proper operators are known. The multitude of operators can get overwhelming, but for most cases, project, match and lookup are enough.

Comparison to Relational

Mongo is very flexible and leaves much of the design open to the context for which the database is built. With this said, it is harder to catch an error than it is usually in SQL. One of the main factors to this is the way references are handled. There isn't anything enforcing relationships, so an error can be introduced very easily, while being hard to detect. It is also easier to fix them than it would be in SQL, due to the flexibility.

The concept of json document makes it easier to design the collections and the relationships between them. Using tables in a SQL Database can seem unfamiliar because it requires a different mindset than the one used in programming. Mongo follows the mentality of object-oriented in comparison to SQL which leans more on set theory and logic.

Advantages and Disadvantages of Mongo in the Exercises

Advantages:

1. Flexibility
2. Easier to understand as it focuses more on steps than nested queries
3. Easier to fix errors
4. Errors don't usually blow you up

Disadvantages:

1. Syntax is overwhelming to write
2. Cases in which aggregates are 0
3. Easier to make errors

4. Errors happen a lot more often
5. Hard to ensure consistency

Conclusion

In conclusion, Mongo seems better in situations in which strictness and consistency are not very important. It presents itself as very good in certain scenarios as IoT and similar fields in which it is hard to ensure that the correct data is sent through.